

Synthetic categorical data generation via variational methods

Natalie Doss

August 31, 2020

1 Introduction

Estimating generative models for complex, high-dimensional data is a common task in machine learning. This task arises in various applications. For instance, synthetic image and document generation have proved important in medical imaging and other areas. In other contexts, small datasets may contain few or none of certain types of examples where prediction accuracy is important. Researchers desire to generate datasets from the same distribution as the original but large enough so these “tail” examples are reasonably frequent.

The difficulty with synthetic data generation in most domains of interest is that the true underlying distribution is possibly high-dimensional and complex, which makes distribution learning computationally infeasible. Bootstrapping, or sampling with replacement, is also not a viable solution when the data are complex and previously unseen examples are desired.

We propose an algorithm for generating synthetic, high-dimensional, categorical data with a potentially complex covariance structure. Since the covariates are categorical, it is not possible to assume the data are Gaussian and estimate the covariance and generate normal data. We demonstrate a method to estimate the covariance for the categorical data and then generate similar categorical data.

2 Model

Consider the following meta-example. Suppose we took a survey of U users, asking them about I categorical items (favorite food, favorite music, favorite car type, and so on). Each user may select one out of K choices per item. We observe data points $x_u \in \mathbb{R}^I$, with each x_{ui} representing the choice of user u in category i . We wish to allow for correlation both within and between covariates. For instance, if a user is “healthy,” perhaps that person is most likely to say their favorite food is apples or carrots, and their favorite activity is exercise or sleeping. This is inspired by the discrete choice models common in the economics literature, e.g., [Braun & McAuliffe \(2010\)](#).

We propose the following latent embedding model. Let the latent variable $z_u \in \mathbb{R}^d$ represents an embedding of user u , and let β_{ik} represent an embedding of category k of item i . Let $\mu \in \mathbb{R}^d, \Sigma \in \mathbb{R}^{d \times d}, \beta_{ik} \in \mathbb{R}^p$. Let η be a function with $\eta : \mathbb{R}^d \rightarrow \mathbb{R}^p$. We assume

$$\begin{aligned} z_u &\sim N(\mu, \Sigma), \\ \pi_{uik} &= \frac{\exp(\eta(z_u)^\top \beta_{ik})}{\sum_{l \in [K]} \exp(\eta(z_u)^\top \beta_{il})}, \\ X_{ui} &\sim \text{Multinom}(\pi_{ui1}, \dots, \pi_{uiK}). \end{aligned} \tag{1}$$

That is, $\mathbb{P}\{X_{ui} = k | z_u, \beta_{ik}\} = \pi_{uik}$. We will typically assume that $d \ll \min(I, K)$. In (1), we have $O(d^2) + O(pIK)$ parameters to estimate; if $p = d$, for example, this is dominated by $O(dIK)$. The observed data are IK -dimensional, so if we estimated the covariance of the original data, we could have $O((IK)^2)$ parameters to estimate, which is far more; (1) allows us to represent our data in a latent space of lower dimension than the original.

In model (1), we might have several z_u vectors that represent “healthy people.” These probably have a large inner product with the β_{ik} ’s associated with favorite food being vegetables and favorite activity being exercise. This model allows us, through the latent z_u variables and β_{ik} parameters, to model our data as having a covariance matrix representing this within- and between- category correlation.

The parameters of the Model (1) are μ, Σ , and the β_{ik} ’s. (And depending on the form we assume for η , we may need to estimate its parameters as well; in practice, we will assume it has a simple linear form.) Given estimates of the parameters, we can generate synthetic data according to (1). We will accordingly estimate these parameters.

Notation We now provide some notation that will be used throughout. Let $\phi_{\mu, \Sigma}$ indicate the density of the $N(\mu, \Sigma)$ distribution. We will use the indices:

$$\begin{aligned} u, v &\in [U] && \text{(user/sample)} \\ i, j &\in [I] && \text{(covariate/item)} \\ k, l &\in [K] && \text{(category)} \end{aligned}$$

For each $u \in [U]$, let $x_{ui} \in \mathbb{R}^K$ be the binary (one-hot) encoding of the categorical variable. Let $x_{1:U}$ indicate the vectors x_1, \dots, x_U , each in \mathbb{R}^{IK} , and let $z_{1:U}$ indicate the vectors z_1, \dots, z_U , each in \mathbb{R}^d . Let $X \in \mathbb{R}^{U \times IK}$ be the matrix whose rows are the x_{ui} . Let $B_i = (\beta_{i1}, \dots, \beta_{iK}) \in \mathbb{R}^{p \times K}$, and let B be collection of matrices B_1, \dots, B_I . For shorthand, let $\theta = (\mu, \Sigma, B)$ represent the model parameters.

Likelihood For a single user u , the joint likelihood is:

$$p(x_u, z_u | \theta) = \left(\prod_{i,k} \pi_{uik}^{x_{uik}} \right) \phi_{\mu, \Sigma}(z_u).$$

And the joint likelihood across U independent users is

$$p(x_{1:U}, z_{1:U} | \theta) = \left(\prod_{u,i,k} \pi_{uik}^{x_{uik}} \right) \left(\prod_u \phi_{\mu, \Sigma}(z_u) \right).$$

Let $b_{ui} = \sum_{l \leq K} \exp(\eta(z_u)^\top \beta_{il})$. Let $f_\theta(z) = \log p_\theta(x_{1:U}, z_{1:U})$. Then

$$f_\theta(z) \stackrel{c}{=} \sum_{u,i,k} x_{uik} \eta(z_u)^\top \beta_{ik} - \sum_{u,i} \log b_{ui} + \sum_u \log \phi_{\mu, \Sigma}(z_u). \quad (2)$$

In the second term in (2), we have used the fact that $\sum_{u,i,k} x_{uik} \log b_{ui} = \sum_{ui} \log b_{ui}$.

3 Variational algorithm

We seek to estimate the parameters in (1). We are particularly interested in Σ , which represents the underlying covariance. A natural approach to optimizing a model containing latent variables is the Expectation Maximization (EM) algorithm. In (2), we cannot compute the E -step in closed form, and so we need to approximate the desired posterior; we will use a variational approximation. For a justification of variational posterior inference in this case, see the tutorial discussion [Blei *et al.* \(2017\)](#) or [Blei & Lafferty \(2007\)](#).

Let $\lambda_u \in \mathbb{R}^d$, and let $V_u \in \mathbb{R}^{d \times d}$ be a diagonal matrix with entries $\nu_{u1}, \dots, \nu_{ud}$. We will make the (mean-field) assumption that the variational posterior q has the form $q(z_{1:U}) = \prod_{u \in [U]} q_u(z_u)$, where $q_u(z_u)$ is the $N(\lambda_u, V_u)$ density. That is,

$$z_u \sim_{i.i.d.} N(\lambda_u, V_u),$$

for $u \in [U]$. The objective is

$$\begin{aligned} \mathcal{L}(\theta) &= \mathbb{E}_q \log p(x, \theta) - \mathbb{E}_q \log q \\ &= \mathbb{E}_{q(z)} f_\theta(z) + \frac{1}{2} \sum_{u \in [U]} \log |V_u|, \end{aligned}$$

where $f_\theta(z)$ is as in (2). We cannot directly calculate $\mathbb{E}_{q(z)} b_{ui}$; this is where the non-conjugacy in this model becomes a problem; see [Blei & Lafferty \(2007\)](#) for further discussion. To handle this, we imitate [Blei & Lafferty \(2007\)](#) and introduce a new set of variational parameters, ζ_u . For any $x, \zeta > 0$,

$$\log x = \log(x/\zeta) + \log \zeta \leq (x/\zeta) - 1 + \log \zeta$$

Using this and Jensen's Inequality,

$$\begin{aligned} \mathbb{E}_q \log b_{ui} &\leq \log(\mathbb{E}_q b_{ui}) \leq \zeta_u^{-1} \mathbb{E}_q b_{ui} - 1 + \log \zeta_u \\ &= \left(\zeta_u^{-1} \sum_{l \leq K} \mathbb{E}_{q(z_u)} \exp(z'_u \beta_{il}) \right) - 1 + \log \zeta_u \\ &= \left(\zeta_u^{-1} \sum_{l \in [K]} \exp(\lambda'_u \beta_{il} + \beta'_{il} V_u \beta_{il} / 2) \right) - 1 + \log \zeta_u \end{aligned}$$

So our objective is

$$\mathcal{L}(\zeta, \lambda_{1:U}, V_{1:U}, B, \mu, \Sigma) \stackrel{c}{=} \sum_{u,i,k} x_{uik} \lambda_u^\top \beta_{ik} - \sum_{u,i,k} \zeta_u^{-1} \exp(\lambda_u^\top \beta_{ik} + \beta_{ik}^\top V_u \beta_{ik}/2) - I \sum_u \log \zeta_u \quad (3)$$

$$- \frac{1}{2} \sum_u ((\lambda_u - \mu)^\top \Sigma^{-1} (\lambda_u - \mu) + \text{tr}(V_u \Sigma^{-1})) \quad (4)$$

$$+ \frac{1}{2} \sum_u (\log |\Sigma|^{-1} + \log |V_u|) \quad (5)$$

E-step of variational EM In the *E*-step, we perform variational posterior inference, i.e., we estimate the variational parameters ζ , $\lambda_{1:U}$, and $V_{1:U}$ in (3). Crucially, we obtain a distinct variational posterior estimate for each $\lambda_u \in \mathbb{R}^d$, $V_u \in \mathbb{R}^{d \times d}$ for $u \in [U]$. In this subsection, I drop the u index on x_u , λ_u , V_u for simplicity. Recall that V has diagonal entries v_1^2, \dots, v_d^2 . I sometimes write this set of entries as a vector $V \in \mathbb{R}^d$.

In optimizing (3), we can obtain $\hat{\zeta}$ in closed form, but not $\hat{\lambda}$, \hat{v}_s^2 . We do a coordinate ascent algorithm, maximizing in ζ , then λ , then ζ again, then V . We repeat until convergence, i.e., until the relative change in the objective (3) is below our desired tolerance level. We maximize by doing the conjugate gradient algorithm for λ and Newton's method in the log space for V . We now provide the necessary calculations to perform these optimizations.

$$\hat{\zeta} = \frac{\sum_{i,k} \exp(\beta_{ik}' \lambda + \beta_{ik}' V \beta_{ik}/2)}{I}$$

The gradients are

$$\nabla_\lambda \mathcal{L}(\lambda) = \sum_{i,k} x_{ik} \beta_{ik} - \zeta^{-1} \sum_{i,k} \beta_{i,k} \exp(\beta_{ik}' \lambda + \beta_{ik}' V \beta_{ik}/2) - \Sigma^{-1}(\lambda - \mu)$$

and

$$\frac{\partial \mathcal{L}(v_s^2)}{\partial v_s^2} = -\zeta^{-1} \sum_{i,k} \frac{\beta_{iks}^2}{2} \exp(\lambda' \beta_{ik} + \beta_{ik}' V \beta_{ik}/2) - \frac{\Sigma_{s,s}^{-1}}{2} + \frac{1}{2v_s^2}$$

And

$$\frac{\partial \mathcal{L}^2(v_s^2)}{\partial (v_s^2)^2} = -\zeta^{-1} \sum_{i,k} \frac{\beta_{iks}^4}{4} \exp(\beta_{ik}' \lambda + \beta_{ik}' V \beta_{ik}/2) - \frac{1}{2v_s^4}$$

We must restrict the v_s^2 's to be positive, so we will do Newton's method in the log space. Recall that Newton's algorithm's updates are:

$$x = x - \frac{f'(x)}{f''(x)}$$

Operating in the log space, let $z = \log x$, so $x = e^z$. So our function $f(x) = f(e^z)$ and we view z as our function argument here.

$$\begin{aligned}\frac{\partial f(e^z)}{\partial z} &= e^z f'(e^z) = x f'(x) \text{ and} \\ \frac{\partial^2 f(e^z)}{\partial^2 z} &= e^{2z} f''(e^z) + e^z f'(e^z) = x^2 f''(x) + x f'(x)\end{aligned}$$

So Newton's algorithm is

$$\log x = \log x - \frac{x f'(x)}{x^2 f''(x) + x f'(x)} = \log x - \frac{f'(x)}{x f''(x) + f'(x)}$$

M -step of variational EM In the M -step, we plug in the approximate posterior $\hat{q} = q_{\hat{\zeta}, \hat{\lambda}_{1:U}, \hat{V}_{1:U}}$ from the E-step and optimize (3) in the parameters μ, Σ, B . It turns out that there is a closed-form solution for these parameters once we have the $\hat{\lambda}_u, \hat{V}_u$'s:

$$\begin{aligned}\hat{\mu} &= \frac{1}{U} \sum_u \lambda_u \\ \hat{\Sigma} &= \frac{1}{U} \left(\sum_u V_u + (\lambda_u - \hat{\mu})(\lambda_u - \hat{\mu})^\top \right)\end{aligned}$$

There is no closed-form solution for $\hat{\beta}_{ik}$. We will do a gradient ascent algorithm and will plug in all parameters that are already estimated. Now

$$\nabla_{\beta_{ik}} L(\beta_{ik}) = \sum_u (x_{uik} \lambda_u - \zeta_u^{-1} (\lambda_u + V_u \beta_{ik}) \exp(\beta'_{ik} \lambda_u + \beta'_{ik} V_u \beta_{ik}/2))$$

The algorithm

Algorithm 1: Parameter estimation via Variational-EM in Model (1). Let $\mathcal{L}(\zeta, \lambda, V, \mu, \Sigma, B)$ be as defined in (3).

Input: Dataset $\{x_u\}_{u \in [U]}$ with each point in \mathbb{R}^{IK} , tolerance level ϵ and stopping time T , initial parameters $\hat{\mu}^{(0)}, \hat{\Sigma}^{(0)}, \hat{B}^{(0)}$.

Output: Estimate $\hat{\mu}, \hat{\Sigma}, \hat{B}$.

While $t \in [T]$ **and likelihood difference ratio is more than** ϵ :

E-step:

Obtain $(\hat{\zeta}^{(t)}, \hat{\lambda}_{1:U}^{(t)}, \hat{V}_{1:U}^{(t)}) \in \operatorname{argmax}_{\zeta, \lambda_{1:U}, V_{1:U}} \mathcal{L}(\zeta, \lambda_{1:U}, V_{1:U}, \hat{\mu}^{(t)}, \hat{\Sigma}^{(t)}, \hat{B}^{(t)})$:

Obtain $\hat{\zeta}^{(t)}$ in closed form ;

Obtain each $\hat{\lambda}_u^{(t)}$ by conjugate gradient ;

Obtain each $\hat{V}_u^{(t)}$ by Newton's method in the log space (for each entry $v_{us}^{(t)}$).

M-step:

Obtain $(\hat{\mu}^{(t+1)}, \hat{\Sigma}^{(t+1)}, \hat{B}^{(t+1)}) \in \operatorname{argmax}_{\mu, \Sigma} \mathcal{L}(\hat{\zeta}^{(t)}, \hat{\lambda}_{1:U}^{(t)}, \hat{V}_{1:U}^{(t)}, \mu, \Sigma, B)$:

$\hat{\mu}_{t+1} = \frac{1}{U} \sum_{u \in [U]} \hat{\lambda}_u^{(t)}$;

$\hat{\Sigma}_{t+1} = \frac{1}{U} \left(\sum_{u \in [U]} \hat{V}_u^{(t)} + \left(\hat{\lambda}_u^{(t)} - \hat{\mu}^{(t+1)} \right) \left(\hat{\lambda}_u^{(t)} - \hat{\mu}^{(t+1)} \right)^\top \right)$.

Report

$$\hat{\mu}^{(T)}, \hat{\Sigma}^{(T)}, \hat{B}^{(T)}. \quad (6)$$

4 Implementation and comparison to other methods

Our model (1) for correlated categorical variables is closely related to the correlated topic model (CTM) of Blei & Lafferty (2007), and we implemented the variational-EM algorithm from Section 5.4.2 by building off the code of Blei & Lafferty (2007) in the C programming language. When we tested our algorithm on data generated from the model (1), we found that it underestimated the values in the covariance matrix Σ , both the variances and covariances. This behavior is not surprising, since the tendency of variational inference to underestimate the posterior covariance may well lead to underestimation of the model covariance in the variational-EM algorithm. To verify that this underestimation is not due to a bug in our adaptation of the CTM code of Blei & Lafferty (2007), we also ran their original code on data generated via a process similar to that of (1). Before explaining this, we describe the relationship between Model (1) and the CTM. We will use the same notation as in Model (1) to be suggestive. We will slightly simplify the correlated topic model to make the correspondences clear; e.g., we will assume each document in a corpus has the same number of words, I .

In the correlated topic model, there are K categories or topics and N words in a dictionary. Let $\mu \in \mathbb{R}^K$, and let $\Sigma \in \mathbb{R}^{K \times K}$ be positive definite. Fix $\alpha_1, \dots, \alpha_K \in \mathbb{R}^N$; each is a discrete probability distribution. We observe a corpus containing U documents, each with I words. The correlated topic model assumes that each word W_i in a document is drawn from

the following generative model.

$$\begin{aligned}
z_{ui} &\sim N(\mu, \Sigma), \\
\pi_{uik} &= \frac{\exp z_{uik}}{\sum_{l \in [K]} \exp z_{uil}}, \\
X_{ui} &\sim_{i.i.d.} \text{Multinom}(\pi_1, \dots, \pi_K), \quad \leftarrow \text{topic assignment} \\
W_{ui} &\sim \text{Multinom}(\alpha_{X_{ui}}). \quad \leftarrow \text{word assignment}
\end{aligned} \tag{7}$$

In (7), the covariance matrix Σ allows for correlation between topics. In (1), we assumed more: that there can be correlation among topics (categories), as well as among the words in a document (covariates/items). We make the following correspondences between CTM and Model (1):

Table 1: Model (1) and CTM Correspondence.

Model (1)	CTM
User/Sample	Document
Covariate/item	Word
Category	Topic

If we modify Model (1) and the CTM in the following ways, we have an exact correspondence. Let $d = p = I * K$, and let Σ be block diagonal with zeros on the inter-covariate blocks. That is, Σ has I^2 block matrices, each $K \times K$, and the diagonal blocks can be nonzero but the off-diagonal blocks are zero. Let η be the identity function, and let $\beta_{ik} = e_{ik}$, the vector with 1 on k th element of group i , and zeros elsewhere. In the CTM, let $N = K$, and let $\alpha_1 = (1, 0, \dots, 0)$, $\alpha_2 = (0, 1, 0, \dots, 0)$, and so on. That is, there is one word per topic; whenever we draw topic k , we are guaranteed that we draw word k . Thus there is a 1-1 correspondence between topics and words; the words are essentially irrelevant. With these simplifications of Model (1) and the CTM, the two models correspond exactly.

We ran the original CTM code, but on data generated according to (7), with the restrictions described above to make the data correspond to (1). That is, we let there be a 1-1 correspondence between word and topic. (However, in our generating process, we did not restrict Σ to be block diagonal.) The CTM correctly estimated the α vectors to be about 1 on one entry and zero on the others. It showed a similar shrinkage to our variational-EM algorithm in the covariance matrix estimation.

Note that when we run CTM, it will estimate one α vector (these are actually called β in the code!) per topic. If there are K topics, the algorithm will estimate $K - 1$ topics so the model is identifiable. For each topic, the α vector is in \mathbb{R}^K . In practice, the α vectors are printed out as one single vector of dimension $(K - 1) * K$. For instance, we tried a number of topics equal to 3, in which case, the CTM algorithm actually estimates that there are 3 topics and 4 terms; this is the identifiability issue. We see in the file “ctm-my-beta,” in the file “final-log-beta.dat,” that we have vectors $(0, 1, 0, 0)$, $(0, 0, 1, 0)$, $(0, 0, 0, 1)$.

We could also estimate Model (1) using a variational autoencoder (VAE). The main differences in what we implemented are that we use mean-field variational inference rather than the amortized mean field inference of VAE. That is, instead of estimating $q_u(z_u)$, the

VAE estimates $q_\phi(z_u)$. See my variational inference sheet for a more complete discussion of VAE's and how they compare to non-amortized variational inference. When we estimated this model using the VAE, we found less shrinkage of the covariance matrix.

We similarly estimated our model using the CTM code from the [Wang & Blei \(2013\)](#) paper. This also implements the CTM model but now uses the Laplace and Delta methods (described therein) to estimate it. These methods showed less shrinkage of the covariance matrix.

5 Alternatives

This section provides some methods that we might also use to estimate this model. We didn't implement these. I use the abbreviate "LLM" for Linear logistic model, for when η is the identity function, and "NLLM" for Non-linear logistic model, for when η is some non linear function, as in a deep neural network. The abbreviation "P" is for parametric, as in amortized mean-field variational inference, and "NP" is for non-parametric, as in the non amortized (classical) mean-field variational inference.

5.1 LLM-NP with Laplace method

The calculations are as in Section 5.2, but now η is the identity function. The following are the gradient and Hessian for a single $z = z_u$; I drop all subscripts u for now. Now

$$\nabla f(z) = \sum_{i,k} x_{ik} \beta_{ik} - \sum_{i,k} \beta_{ik} \pi(z, \beta_{ik}) - \Sigma^{-1}(z - \mu)$$

And

$$\nabla^2 f(z)_{s,t} = - \sum_{i,k} \beta_{iks} \pi(z, \beta_{ik}) \left(\beta_{ikt} - \sum_{l \in [K]} \beta_{ilt} \pi(z, \beta_{il}) \right) - \Sigma_{s,t}^{-1}$$

And once we have q , we know that our objective is as follows. I let $\hat{\lambda}_u := \hat{\lambda}(x_u)$ and similarly for \hat{V}_u . I use \hat{q} to indicate $\hat{q}(x_1), \dots, \hat{q}(x_U)$. Let $\xi_u \sim_{i.i.d.} N(0, I_d)$. Using a single sample to approximate $\mathbb{E}_q \log b_{ui}$,

$$\begin{aligned} \mathcal{L}(\hat{q}) &\approx \sum_{u,i,k} x_{uik} \hat{\lambda}'_u \beta_{ik} - \sum_{u,i} \mathbb{E}_{\hat{q}} \sum_{l \in [K]} \exp \left((\hat{\lambda}_u + \hat{V}_u^{1/2} \xi_u)' \beta_{il} \right) - \\ &\quad \frac{1}{2} \sum_u \left((\hat{\lambda}_u - \mu)' \Sigma^{-1} (\hat{\lambda}_u - \mu) + \text{tr}(V \Sigma^{-1}) + \log |V| \right) \end{aligned}$$

E-step: For each $u \in [U]$, $\hat{q}(z_u)$ is the $N(\hat{\lambda}_u, \hat{V}_u)$ density, where

$$\begin{aligned} \hat{\lambda}_u &= \hat{z}_u = \operatorname{argmax} f(z_u) \\ \hat{V}_u &= -\nabla^2 f(\hat{z}_u)^{-1} \end{aligned}$$

These are both found using the gradient and Hessian, calculated above.

M-step: Using the objective after finding \hat{q} , we see that:

$$\begin{aligned}\hat{\mu} &= \frac{1}{U} \sum_u \hat{\lambda}_u \\ \hat{\Sigma} &= \frac{1}{U} \sum_u (\hat{\lambda}_u - \hat{\mu})(\hat{\lambda}_u - \hat{\mu})' + \frac{1}{U} \sum_u \hat{V}_u\end{aligned}$$

For β_{ik} , we don't have an analytic solution, but we can do gradient ascent. The gradient is:

$$\begin{aligned}\nabla_{\beta_{ik}} \mathcal{L} &= \sum_u x_{uik} \hat{\lambda}_u - \sum_u \frac{(\hat{\lambda}_u + \hat{V}_u^{1/2} \xi_u) \exp\left((\hat{\lambda}_u + \hat{V}_u^{1/2} \xi_u)' \beta_{ik}\right)}{\sum_{l \leq K} \exp\left((\hat{\lambda}_u + \hat{V}_u^{1/2} \xi_u)' \beta_{il}\right)} \\ &= \sum_u x_{uik} \hat{\lambda}_u - \sum_u \hat{a}_u \pi(\hat{a}_u, \beta_{ik})\end{aligned}$$

where $\hat{a}_u = \hat{\lambda}_u + \hat{V}_u^{1/2} \xi_u$.

If we have the identity β , everything is just as in Section 5.1, except now the gradients are as follows. Note how this matches the CTM calculations of Wang & Blei (2013) for the latent variable in that model. That is, replace their $t(z)$ with our $\sum_i x_i$ where $x_i \in \mathbb{R}^K$; everything is just the same.

$$\begin{aligned}\nabla f(z) &= \sum_i x_i - I\pi - \Sigma^{-1}(z - \mu) \\ \nabla^2 f(z)_{st} &= \pi_s(\mathbf{1}\{s = t\} - \pi_t) - \Sigma_{st}^{-1}\end{aligned}$$

Now we don't have B in the model anymore; we just have μ, Σ . And their updates will be as in Section 5.1.

5.2 NLLM-NP with Laplace method

Let $J(\eta)$ be the Jacobian of η ; note $J(\eta) \in \mathbb{R}^{p \times d}$. And let

$$H(\eta) = (H(\eta_1), \dots, H(\eta_p))$$

be a tensor that is the array of the Hessians of the components of η . So we write

$$H(\eta)'_{s,t} \beta := \sum_{j \leq p} \frac{\partial^2 \eta_j(z)}{\partial z_s \partial z_t} \beta_j$$

I sometimes abbreviate $J(\eta), H(\eta)$ to just J, H . And I sometimes drop the $\eta(z)$ and just write η . Now

$$\nabla f(z) = \sum_{i,k} x_{ik} J(\eta)' \beta_{ik} - \sum_{ik} \frac{J(\eta)' \beta_{ik} \exp(\eta' \beta_{ik})}{b_i} - \Sigma^{-1}(z - \mu)$$

We cannot find \hat{z} in closed form, but we can do gradient ascent or some other algorithm to find it. And for the Hessian, first recall that for any function $h(z)$,

$$\frac{\partial^2(\log h(z))}{\partial z^2} = \frac{h''(z)}{h(z)} - \left(\frac{h'(z)}{h(z)} \right)^2$$

Now

$$\begin{aligned} \nabla^2 f(z)_{s,t} &= \sum_{i,k} x_{ik} H(\eta)'_{s,t} \beta_{ik} \\ &\quad - \sum_i \left(\sum_k \frac{\exp(\eta' \beta_k) \left(H'_{s,t} \beta_k + (J'_{[s]} \beta_k * J'_{[t]} \beta_k) \right)}{b_i} - \frac{\sum_k J'_{[s]} \beta_k \exp(\eta' \beta_k) \sum_l J'_{[t]} \beta_l \exp(\eta' \beta_l)}{b_i^2} \right) \\ &\quad - \Sigma_{s,t}^{-1} \end{aligned}$$

Now if we do a variational-EM algorithm, the M -step will involve taking derivatives of $f_{\theta,x}(z)$ with respect to the parameters θ . If we have the $N(\mu, \Sigma)$ prior, then $\theta_1 = (\mu, \Sigma)$, and the updates are the sufficient statistics as in Section 3. For $\tilde{\theta}_2$, we will need the derivatives of η with respect to these parameters. And as in the discussion in the variational inference sheet, we can approximate the integral via sampling, since we have from the E -step the $\hat{\lambda}, \hat{V}$ for q . Suppose we approximate the integral via one sample ξ . Write $\tilde{\eta}_u = \eta(\hat{\lambda}(x_u) + \hat{V}^{1/2}(x_u)\xi_u)$.

$$\nabla_{\beta_{ik}} \mathcal{L} = \sum_u x_{uik} \tilde{\eta}_u - \sum_u \frac{\tilde{\eta}_u \exp(\tilde{\eta}'_u \beta_{ik})}{\sum_l \exp(\tilde{\eta}'_u \beta_{il})}$$

5.3 LLM-NP with sampling

5.3.1 Objective (ELBO)

Now instead of introducing ζ to compute $\mathbb{E}_{q(z_u)} \log b_{ui}$, we do the following.

$$\begin{aligned} \mathbb{E}_{q(z_u)} \log \sum_{k \leq K} \exp(z'_u \beta_{ik}) &= \mathbb{E}_{\xi_u \sim N(0, I_d)} \log \sum_{k \leq K} \exp(\lambda'_u \beta_{ik} + \xi'_u V_u^{1/2} \beta_{ik}) \\ &\approx \log \sum_{k \leq K} \exp(\lambda'_u \beta_{ik} + \xi'_u V_u^{1/2} \beta_{ik}) \end{aligned}$$

where $\xi_u \sim N(0, I_d)$. That is, I'm approximating the integral with a single draw from the distribution. We could use more draws to get a better approximate. Our full ELBO now is:

$$\begin{aligned} \mathcal{L}(\lambda_{1:U}, V_{1:U}, B, \mu, \Sigma) &= \sum_{u,i,k} x_{uik} \lambda'_u \beta_{ik} - \sum_{u,i} \log \sum_{k \leq K} \exp(\lambda'_u \beta_{ik} + \xi'_u V_u^{1/2} \beta_{ik}) \\ &\quad + U \log |\Sigma^{-1}| - \frac{1}{2} \sum_u ((\lambda_u - \mu)' \Sigma^{-1} (\lambda_u - \mu) + \text{tr}(V_u^{1/2} \Sigma^{-1} V_u^{1/2})) \\ &\quad + \frac{\sum_u \log |V_u|}{2} \end{aligned}$$

5.3.2 Variational algorithm

The gradients are (dropping indices for now):

$$\begin{aligned}\nabla_{\lambda}\mathcal{L}(\lambda) &= \sum_{i,k} x_{ik}\beta_{ik} - \sum_{i \leq I} \frac{\sum_{k \leq K} \beta_{ik} \exp(\lambda' \beta_{ik} + \xi' V^{1/2} \beta_{ik})}{\sum_{l \leq K} \exp(\lambda' \beta_{il} + \xi' V^{1/2} \beta_{ik})} - \frac{1}{2} \Sigma^{-1}(\lambda - \mu) \\ \frac{\partial \mathcal{L}(v_s^2)}{\partial v_s^2} &= - \sum_{i \leq I} \frac{\sum_{k \leq K} \frac{\xi_s \beta_{iks}}{2v_s} \exp(\lambda' \beta_{ik} + \xi' V^{1/2} \beta_{ik})}{\sum_{l \leq K} \exp(\lambda' \beta_{il} + \xi' V^{1/2} \beta_{ik})} - \frac{\Sigma_{s,s}^{-1}}{2} + \frac{1}{2v_s^2} \\ &= - \frac{\xi_s}{2v_s} \sum_{i \leq I} \frac{\sum_{k \leq K} \beta_{iks} \exp(\lambda' \beta_{ik} + \xi' V^{1/2} \beta_{ik})}{\sum_{l \leq K} \exp(\lambda' \beta_{il} + \xi' V^{1/2} \beta_{ik})} - \frac{\Sigma_{s,s}^{-1}}{2} + \frac{1}{2v_s^2}\end{aligned}$$

For the M -step, the solutions for μ, Σ are the same as in previous sections. And

$$\nabla_{\beta_{ik}} \mathcal{L}(\beta_{ik}) = \sum_u x_{uik} \lambda_u - \sum_u \frac{(\lambda_u + V_u^{1/2} \xi_u) \exp(\lambda'_u \beta_{ik} + \xi'_u V_u^{1/2} \beta_{ik})}{\sum_{l \leq K} \exp(\lambda'_u \beta_{il} + \xi'_u V_u^{1/2} \beta_{ik})}$$

5.4 LLM-NP with linear approximation and β prior

Let $z_u, \beta_{ik} \in \mathbb{R}^d$. And let $\nu, \mu \in \mathbb{R}^d$ and $\Omega, \Sigma \in \mathbb{R}^{d \times d}$. Now we place a prior on β ; here is the data-generating process.

$$z_u \sim N(\mu, \Sigma) \tag{8}$$

$$\beta_{ik} \sim N(0, \gamma^2 I_d) \tag{9}$$

$$\mathbb{P}\{X_{ui} = k | z_u, \beta_{ik}\} = \frac{\exp(z'_u \beta_{ik})}{\sum_{l \in [K]} \exp(z'_u \beta_{il})} \tag{10}$$

For variational inference, we use the families:

$$q(z_u) = N(\lambda_u, V_u)$$

$$q(\beta_{ik}) = N(\psi_{ik}, W_{ik})$$

where W_u, V_{ik} are all diagonal matrices with entries v_{iks}, w_{us} for $s \in [d]$. We find the parameters of q to maximize

$$\mathbb{E}_{z \sim q} \log p(x, \theta, B) - \mathbb{E}_q \log q$$

5.4.1 Objective (ELBO)

The joint across U independent users is

$$p(x_{1:U}, z_{1:U}, B) = \left(\prod_{u,i,k} \pi(z_u, \beta_{ik})^{x_{uik}} \right) \left(\prod_u \phi_{\mu, \Sigma}(z_u) \right) \left(\prod_{i,k} \phi_{0, \gamma^2 I_d}(\beta_{ik}) \right)$$

Now let $b_{ui} = \sum_{l \leq K} \exp(z'_u \beta_{il})$. We have

$$\log p(x, \theta, B) \stackrel{c}{=} \sum_{u,i,k} x_{uik} z'_u \beta_{ik} - \sum_{u,i} \log b_{ui} + \sum_u \log \phi_{\mu, \Sigma}(z_u) + \sum_{i,k} \log \phi_{0, \gamma^2 I_d}(\beta_{ik})$$

Note that in the second term, we used the fact that $\sum_{u,i,k} x_{u,i,k} b_{ui} = \sum_{ui} b_{ui}$. Now to help handle the expectation of this term, we introduce a new set of variational parameters, ζ_u . Note that for any $x, \zeta > 0$,

$$\log x = \log(x/\zeta) + \log \zeta \leq (x/\zeta) - 1 + \log \zeta$$

Using this and Jensen,

$$\begin{aligned} \mathbb{E}_q \log b_{ui} &\leq \log(\mathbb{E}_q b_{ui}) \leq \zeta_u^{-1} \mathbb{E}_q b_{ui} - 1 + \log \zeta_u \\ &= \left(\zeta_u^{-1} \sum_{l \leq K} \mathbb{E}_q \exp(z'_u \beta_{il}) \right) - 1 + \log \zeta_u \end{aligned}$$

We have by the diagonality of the variance matrices and the independence of θ_u, β_{ik} ,

$$\mathbb{E}_{q(\beta_{ik})} \mathbb{E}_{q(z_u)} e^{z'_u \beta_{il}} = \prod_{s \in [d]} f_{uils}$$

where $f_{uils} = \mathbb{E}_{q(z_u)} \mathbb{E}_{q(\beta_{ik})} \exp(z_u[s] \beta_{il}[s])$. See Section 5.4.3 for its full form.

$$\mathbb{E}_q \log b_{ui} \leq \left(\zeta_u^{-1} \sum_{l \in [K]} \prod_{s \in [d]} f_{uils} \right) - 1 + \log \zeta_u$$

And

$$\begin{aligned} H(q(\theta_u)) &= \frac{-\log|V_u^{-1}|}{2} = \frac{\log|V_u|}{2} \\ H(q(\beta_{ik})) &= \frac{-\log|W_{ik}^{-1}|}{2} = \frac{\log|W_{ik}|}{2} \end{aligned}$$

So our full ELBO is:

$$\begin{aligned} &\sum_{u,i,k} x_{uik} \lambda'_u \lambda_{ik} - \left(\sum_{u,i} \zeta_u^{-1} \sum_{l \in [K]} \prod_{s \in [d]} f_{uils} \right) - \sum_{u,i} \log \zeta_u + U \log |\Sigma^{-1}| + IKd \log \frac{1}{\gamma^2} \\ &- \frac{1}{2} \sum_u ((\lambda_u - \mu)' \Sigma^{-1} (\lambda_u - \mu) + \text{tr}(V_u^{1/2} \Sigma^{-1} V_u^{1/2})) - \frac{1}{2} \sum_{i,k} \left(\frac{\|\psi_{ik}\|_2^2}{\gamma^2} + \frac{\text{tr}(W_{ik})}{\gamma^2} \right) \\ &+ \frac{\sum_u \log|V_u| + \sum_{ik} \log|W_{ik}|}{2} \end{aligned}$$

5.4.2 Variational algorithm

Variational E-step updates I'm letting v_u, w_{ik} be the vectors in question for the diagonal matrices. The elements are v_{us} or $v_u[s]$, either way.

$$\partial_{\lambda_{ut}} L(\lambda_{ut}) = \sum_{i,k} x_{uik} \psi_{ik}[t] - \Sigma^{-1}(\lambda_u - \mu)[t] - \sum_{i,k} \frac{\partial f(\lambda_{ut}, v_{ut}^2, \psi_{ikt}, w_{ikt}^2)}{\partial \lambda_{ut}} \prod_{s \neq t} f(\lambda_{us}, v_{us}^2, \psi_{iks}, w_{iks}^2)$$

$$\partial_{\psi_{ikt}} L(\psi_{ikt}) = \sum_u x_{uik} \lambda_u[t] - \frac{1}{\gamma^2} \psi_{ik}[t] - \sum_u \frac{\partial f(\lambda_{ut}, v_{ut}^2, \psi_{ikt}, w_{ikt}^2)}{\partial \psi_{ikt}} \prod_{s \neq t} f(\lambda_{us}, v_{us}^2, \psi_{iks}, w_{iks}^2)$$

$$\partial_{v_{ut}} L(v_{ut}^2) = -\frac{\text{diag}(\Sigma^{-1})[t]}{2} + \frac{1}{2v_u^2[t]} - \sum_{i,k} \frac{\partial f(\lambda_{ut}, v_{ut}^2, \psi_{ikt}, w_{ikt}^2)}{\partial v_{ut}^2} \prod_{s \neq t} f(\lambda_{us}, v_{us}^2, \psi_{iks}, w_{iks}^2)$$

$$\partial_{w_{ikt}} L(w_{ikt}^2) = -\frac{1}{2} + \frac{1}{2w_{ik}^2[t]} - \sum_u \frac{\partial f(\lambda_{ut}, v_{ut}^2, \psi_{ikt}, w_{ikt}^2)}{\partial w_{ikt}^2} \prod_{s \neq t} f(\lambda_{us}, v_{us}^2, \psi_{iks}, w_{iks}^2)$$

And in closed form, for each $u \in [U]$,

$$\hat{\zeta}_u = \frac{-\sum_{i,k} \prod_{s \leq d} f_{uiks}}{I}$$

Sufficient statistics for M step

$$\begin{aligned} \hat{\mu} &= \frac{1}{U} \sum_u \lambda_u \\ \hat{\Sigma} &= \frac{1}{U} \sum_u (\lambda_u - \hat{\mu})(\lambda_u - \hat{\mu})' + \frac{1}{U} \sum_u V_u \\ \hat{\gamma}^2 &= \frac{1}{IKd} \sum_{i,k} (\|\psi_{ik}\|_2^2 + \text{tr}(W_{ik})) \end{aligned}$$

5.4.3 MGF calculation

To evaluate $\mathbb{E}_{z_u \sim q} \mathbb{E}_{\beta_{ik} \sim q} \exp(z'_u \beta_{ik})$, first consider the following. Let

$$\begin{aligned} z &\sim N(\lambda, v^2) \\ \beta &\sim N(\psi, w^2) \end{aligned}$$

Then

$$\begin{aligned}
\mathbb{E}_\beta \mathbb{E}_z \exp(z\beta) &= \mathbb{E}_\beta \exp(\lambda\beta + v^2\beta^2/2) \\
&= \mathbb{E}_\beta \exp\left(\frac{v^2}{2} \left(\beta^2 + \frac{2\lambda}{v^2}\beta + \frac{\lambda^2}{v^4} - \frac{\lambda^2}{v^4}\right)\right) \\
&= \mathbb{E}_\beta \exp\left(\frac{v^2}{2} \left(\beta + \frac{\lambda}{v^2}\right)^2 - \frac{\lambda^2}{2v^2}\right) \\
&= e^{-\lambda^2/2v^2} \mathbb{E}_\beta \exp\left(\frac{v^2}{2} \left(\beta + \frac{\lambda}{v^2}\right)^2\right)
\end{aligned}$$

Now $\beta + \lambda/v^2 \sim N(\psi + \lambda/v^2, w^2)$, so

$$\begin{aligned}
\beta + \frac{\lambda}{v^2} &\sim wN\left(\frac{\psi v^2 + \lambda}{wv^2}, 1\right) \Rightarrow \\
\left(\beta + \frac{\lambda}{v^2}\right)^2 &\sim w^2 \chi_1^2\left(\left(\frac{\psi v^2 + \lambda}{wv^2}\right)^2\right) \\
&\sim \frac{1}{v^4} \chi_1^2(\psi^2 v^4 + \lambda^2 + 2\psi\lambda v^2)
\end{aligned}$$

Using the moment-generating function for the non-central chi square distribution, we have:

$$\begin{aligned}
f(\lambda, v^2, \psi, w^2) &= \mathbb{E}_\beta \mathbb{E}_z \exp(z\beta) = \frac{1}{\sqrt{1 - v^2 w^2}} \exp\left(-\frac{\lambda^2}{2v^2}\right) \exp\left(\frac{v^2}{2v^4} \frac{\psi^2 v^4 + \lambda^2 + 2\psi\lambda v^2}{1 - v^2 w^2}\right) \\
&= \frac{1}{\sqrt{1 - v^2 w^2}} \exp\left(\frac{\psi^2 v^4 + \lambda^2 + 2\psi\lambda v^2 - \lambda^2 + \lambda^2 v^2 w^2}{2v^2(1 - v^2 w^2)}\right) \\
&= \frac{1}{\sqrt{1 - v^2 w^2}} \exp\left(\frac{\psi^2 v^2 + 2\psi\lambda + \lambda^2 w^2}{2(1 - v^2 w^2)}\right)
\end{aligned}$$

Let $g(\lambda, v^2, \psi, w^2) = \frac{\psi^2 v^2 + 2\psi\lambda + \lambda^2 w^2}{2(1 - v^2 w^2)}$. The gradients of this with respect to each parameter are:

$$\begin{aligned}
\frac{\partial f}{\partial \lambda} &= \frac{2\lambda w^2 + 2\psi}{2(1 - v^2 w^2)^{3/2}} \exp(g(\lambda, v^2, \psi, w^2)) \\
\frac{\partial f}{\partial \psi} &= \frac{2\psi v^2 + 2\lambda}{2(1 - v^2 w^2)^{3/2}} \exp(g(\lambda, v^2, \psi, w^2))
\end{aligned}$$

and

$$\begin{aligned}
\frac{\partial f}{\partial v^2} &= \frac{\partial g / \partial v^2}{2(1 - v^2 w^2)^{1/2}} \exp(g(\lambda, v^2, \psi, w^2)) + \frac{w^2}{2(1 - v^2 w^2)^{3/2}} \exp(g(\lambda, v^2, \psi, w^2)) \\
\frac{\partial f}{\partial w^2} &= \frac{\partial g / \partial w^2}{2(1 - v^2 w^2)^{1/2}} \exp(g(\lambda, v^2, \psi, w^2)) + \frac{v^2}{2(1 - v^2 w^2)^{3/2}} \exp(g(\lambda, v^2, \psi, w^2))
\end{aligned}$$

where

$$\frac{\partial g}{\partial v^2} = \frac{(1 - v^2 w^2) \lambda^2 + (\psi^2 v^2 + 2\psi \lambda + \lambda^2 w^2) w^2}{2(1 - v^2 w^2)^2}$$

$$\frac{\partial g}{\partial w^2} = \frac{(1 - v^2 w^2) \lambda^2 + (\psi^2 v^2 + 2\psi \lambda + \lambda^2 w^2) v^2}{2(1 - v^2 w^2)^2}$$

References

- Blei, David M., & Lafferty, John D. 2007. A correlated topic model of science. *Annals of Applied Statistics*, **1**(1), 17–35.
- Blei, David M., Kucukelbir, Alp, & McAuliffe, Jon D. 2017. Variational Inference: A Review for Statisticians. *Journal of the American Statistical Association*, **112**(518), 859–877.
- Braun, Michael, & McAuliffe, Jon. 2010. Variational Inference for Large-Scale Models of Discrete Choice. *Journal of the American Statistical Association*, **105**(489), 324–335.
- Wang, Chong, & Blei, David M. 2013. Variational inference in nonconjugate models. *Journal of Machine Learning Research*, **14**, 1005–1031.